

ChrisBuen

AI Automation Builder

I build AI-assisted systems that connect research, operational rules, document production, and human approval. My strongest work sits where a normal prompt is not enough: multi-step workflows, source-bound research, controlled document generation, reusable agent instructions, and evidence that a reviewer can inspect.

GitHub: github.com/ChrisBuen

Evidence at a glance

Public evidence	Result
Project folders	20
Runnable unit tests	23 passed
Converted knowledge files	3,072
Advice batch	12 work items / 24 DOCX files
Contract classification set	12 cases / 1.000 accuracy
Largest recorded n8n file fetch	24,112,730 bytes

Three projects to inspect first

SharePoint file delivery through n8n. An AI-facing tool searched SharePoint through Microsoft Graph, validated the selected file, enforced a 25 MiB response limit, and returned binary data or a short-lived link. The public package includes a sanitized n8n export, an API deployment helper, fictional request/response fixtures, architecture, and three tests.

Governed contract automation. A precedent-led system classified contract families, retrieved clauses, assembled review-ready Word documents, recorded provenance, and applied zero-tolerance leakage and invented-fact gates. A fixed private evaluation recorded 0.917 retrieval precision at 3, full provenance coverage, and zero cross-matter leakage. Operational reliance remained blocked.

Advice-document production. Structured records were transformed into conditional narratives and governed DOCX outputs while preserving styles, tables, headers, footers, and images. A staged batch generated 24 Word files across 12 work items and passed recorded QA.

What I bring to an automation role

Workflow engineering

- n8n workflow design, webhook contracts, Microsoft Graph calls, binary file handling, and API deployment;
- explicit input validation, error paths, idempotent updates, and environment-variable configuration;
- filesystem and JSON contracts for handoffs between agents, repositories, and workflows.

Research and decision systems

- source-bound company, infrastructure, and due-diligence research routines;
- evidence matrices, source logs, confidence labels, and evidence-gap outputs;
- specialist decision councils that keep dissent, risks, conditions, and human ownership visible.

Documents, reports, and knowledge

- conditional DOCX generation and package-level QA;
- Markdown conversion and provenance indexes across a 3,072-file knowledge collection;
- recurring chart/report workflows, report archetypes, and DOCX/PDF release checks.

Product and content operations

- modular prompt libraries, transcript-to-script pipelines, voice/podcast tooling, and SVG infographic skills;
- TypeScript contribution-scoring logic for a running-club MVP;
- responsive AI education site work and a prioritized technical/content SEO audit.

Evidence standard

The portfolio distinguishes implemented/evaluated systems, runnable public proof, source extracts, and design-stage scaffolds. Fictional fixtures are labeled. Private client data, contract language, credentials, internal URLs, and regulated case material are excluded. The recorded numbers are implementation evidence, not ROI or adoption claims.

Start here

- [README.md](#) — one-minute portfolio scan
- [PROJECTS.md](#) — all 20 project groups
- [projects/n8n-deployment-and-tools/](#)
- [projects/governed-contract-automation/](#)
- [projects/advice-document-pipeline/](#)
- [downloads/chrisbuen-ai-automation-casebook.pdf](#)

Contact: <https://github.com/ChrisBuen>