

ChrisBuen — Project Evidence Appendix

AI Automation Builder

This appendix is the file-level index for 20 project groups. It distinguishes four evidence types:

- **Implemented / evaluated:** working private system with a safe public extract or aggregate.
- **Runnable public proof:** public fixture and code can be executed locally.
- **Source extract:** authored code, captures, or outputs are published; full deployment remains private.
- **Design-stage / scaffold:** included honestly and not presented as completed.

The public repository contains 23 passing Python unit tests plus one recorded Vitest test from the source Club 4301 project. Private provenance, hashes, and sanitization records remain ignored.

Contact: github.com/ChrisBuen

01 — n8n deployment and SharePoint tools

Status: implemented; sanitized public extraction.

Proof files

- `projects/n8n-deployment-and-tools/workflows/sharepoint-file-delivery.workflow.json`
- `projects/n8n-deployment-and-tools/src/deploy_workflow.py`
- `projects/n8n-deployment-and-tools/samples/request.json`
- `projects/n8n-deployment-and-tools/samples/response.json`
- `projects/n8n-deployment-and-tools/tests/test_deploy_workflow.py`
- `artifacts/diagrams/n8n-sharepoint-delivery.png`

Recorded result: 24,112,730-byte audio fetch, 25 MiB guard, temporary PDF delivery, extension correction, idempotent redeployment.

Boundary: production IDs, URLs, OAuth references, and credentials are replaced.

02 — Governed contract automation

Status: implemented and privately evaluated.

Proof files: runnable metric gates, three tests, synthetic cases, sanitized aggregate, control diagram.

Recorded result: 1.000 classification accuracy across 12 cases; 0.917 retrieval precision at 3; 1.000 provenance; zero cross-matter leakage; operational reliance false.

Boundary: no precedent language, party details, or generated client agreement.

03 — Advice-document pipeline

Status: implemented; private batch evidence.

Proof files: `document_qa.py`, three tests, fictional input/output, test receipt, pipeline diagram.

Recorded result: 12 work items, 24 DOCX outputs, recorded batch QA pass, preserved styles/tables/headers/footers/images in v11.

Boundary: fictional public sample; no advice or client file.

04 — Applicant assessment architecture

Status: runnable public control extract.

Proof files: `stage_gate.py`, three tests, fictional case, evidence-gap handoff, state diagram.

What it proves: missing evidence blocks progression and automation cannot approve or decline.

Boundary: no policy thresholds, forms, regulated identifiers, or applicant data.

05 — AI-readable knowledge base

Status: implemented; runnable public fixture.

Proof files: original-derived conversion and index scripts, source manifest, converted Markdown, master index, collection chart.

Recorded result: 3,072 private files converted across research, marketing, education, compliance, client, and fund collections.

Boundary: category counts only; corpus and private metadata excluded.

06 — Research and due-diligence skill library

Status: implemented privately; neutralized public extracts.

Proof files: company research, infrastructure readiness, and DD orchestration skill extracts; fictional input and evidence-gap output.

What it proves: primary-source hierarchy, discrepancy handling, confidence, routing, and human-review output contracts.

Boundary: company policy, applicants, staff contacts, and private tool routes removed.

07 — Decision council

Status: runnable public proof.

Proof files: decision schema, validator, fictional brief, two-seat output, two tests, architecture diagram.

What it proves: multiple perspectives, evidence used, risk, conditions, confidence, dissent, and human ownership.

Boundary: fictional decision; no investment recommendation.

08 — Market-report automation

Status: runnable public proof with synthetic data.

Proof files: CSV validator/SVG renderer, six-observation fixture, source manifest, generated chart, two tests.

Private components recorded: market analysis, unlock and ETF-flow collectors, branded charts, DOCX assembly.

Boundary: public values are invented; no market conclusion.

09 — Research report builder

Status: runnable staged workspace.

Proof files: five original-derived scripts plus generated research brief, source log, evidence matrix, and briefing-note skeleton.

What it proves: research planning precedes drafting and supports nine report archetypes.

Boundary: demo workspace contains no external claim.

10 — Learning-platform publishing

Status: scaffold only.

Proof files: complete safe source snapshot and a fictional proposed publication contract.

What was found: one project guidance file stating the codebase was new; no publisher, cron, payload, or test record.

Boundary: not promoted as an implemented LearnWorlds integration.

11 — Voice and podcast renderer

Status: working source extract.

Proof files: original parser and audio stitcher, fictional script, limitations matrix, three parser tests.

What it proves: speaker contract, chunking, sample-rate checks, mono conversion, pause insertion, peak normalization, PCM export.

Boundary: no voice references, imitation audio, or model weights.

12 — SVG infographic skill

Status: complete source extract.

Proof files: full [SKILL.md](#), source syntax, editable SVG, rendered WEBP, PowerShell wrapper.

What it proves: a reusable agent skill shipped with instructions and a concrete visual artifact.

Boundary: third-party dependencies omitted; upstream licenses apply.

13 — Research-to-content pipeline

Status: runnable public proof.

Proof files: content-package builder, fictional transcript, generated JSON package, two tests, architecture diagram.

What it proves: timestamped evidence flows into a constrained brief, script beat, and storyboard scene.

Boundary: no third-party transcript or private channel strategy.

14 — Club 4301 MVP

Status: working source-project extract.

Proof files: original `scoring.ts`, `scoring.test.ts`, `recap.ts`, and `types.ts`; source test receipt.

Recorded result: one Vitest file and one test passed on August 11, 2026.

Boundary: full app and deployment configuration await a separate review.

15 — AI content and SEO system

Status: implemented site/workspace evidence.

Proof files: final local desktop and mobile homepage captures; sanitized issue-to-action audit.

Recorded audit evidence: roughly 23 live article records, 11 without tracked incoming links, four missing hub articles, and specific priority actions.

Boundary: no ranking, traffic, revenue, or post-fix outcome claim.

16 — Prompt decomposition library

Status: original-derived public rewrite.

Proof files: shared policy; quiz, flashcard, and FAQ modules; fictional source and outputs.

What it proves: one responsibility and one output contract per module, with source-bound generation.

Boundary: course content, workflow IDs, and proprietary prompts removed.

17 — Document composition system

Status: implemented through the portfolio document set.

Proof files: neutral design tokens, six page archetypes, before/after diagram, editable DOCX and fixed PDF outputs.

What it proves: content hierarchy, evidence labels, charts, metadata, and render QA are treated as system controls.

Boundary: private brand assets and templates excluded.

18 — SEO automation architecture

Status: design-stage.

Proof files: workflow contract, fictional keyword input, diagram, promotion checklist.

What it proves: the planned stages and stop conditions are explicit.

Boundary: no runnable n8n export or publication receipt; therefore no implementation claim.

19 — GitHub-to-n8n handoff

Status: runnable public validator.

Proof files: directory contract, validator, complete fictional case, two tests.

What it proves: required file checks, case-ID reconciliation, allowed status values, and mandatory human review.

Boundary: no repository, branch, webhook, workflow, or messaging identifiers.

20 — Annual-report workspace

Status: staged integration design.

Proof files: workspace manifest, source-to-output checklist, assembly diagram, links to reused report/chart/document packages.

What it proves: component reuse and output ownership across a long-form reporting workspace.

Boundary: no private annual report or duplicated internal asset.

Verification index

Runnable Python tests

Package	Tests
Advice-document pipeline	3
Assessment automation	3
Decision council	2
GitHub-to-n8n handoff	2
Governed contract automation	3
Market-report automation	2
n8n deployment and tools	3
Research-to-content pipeline	2
Voice/podcast parser	3
Total	23

Release checks required before GitHub publication

- scan public text/code/JSON for credentials, private emails, local paths, client names, and live infrastructure;
- inspect image and office-document metadata;
- validate JSON and SVG syntax;
- validate DOCX packages and render every document to PDF;
- inspect page images for clipping, blank pages, weak density, and inconsistent hierarchy;
- verify relative links and expected download filenames;
- confirm **private/** is ignored and absent from the public file list;
- obtain user approval before any GitHub push.

Contact

<https://github.com/ChrisBuen>